

# Homework 1 Tips

**Week 2 – Part 2**

**Amittai Aviram – 15 September 2020**

# Python Resources

## Threads and Processes

- Threads
  - `import threading`
  - User-level threads—scheduled by the Python runtime
  - Do not map onto the machine's CPU cores—all threads are executed on a single core
  - Share the processes memory—global variables, etc.
  - Enable the process to get work done while one or another thread waits for resources
  - Best for *network-bound* applications
- Processes
  - `import multiprocessing`
  - Kernel-level processes—scheduled by the OS
  - The OS can schedule them on available CPU cores (like other processes)
  - Shared data must be communicated between processes
  - The OS incurs the cost of setting up *context* for each new process
  - Enable tasks to be performed *at the same time*
  - Best for *compute-bound* applications

# Similar API

## The Basics

```
from threading import Thread
from typing import List
```

```
def foo(x: int, y: int, z: List[int]):
    z.append(x * y)
    z.append(x + y)
```

```
if __name__ == '__main__':
    result = []
    thread_0 = Thread(target=foo, args=(17, 19, result))
    thread_0.start()
    thread_0.join()
    print(result)
```

```
from multiprocessing import Array, Process
```

```
def foo(x: int, y: int, z: Array):
    z[0] = (x * y)
    z[1] = (x + y)
```

```
if __name__ == '__main__':
    result = Array('i', 2)
    process_0 = Process(target=foo, args=(17, 19, result))
    process_0.start()
    process_0.join()
    print(result[:])
```


# Synchronization and Communication

See Python Documentation

- Communication
  - Queue (both threading and multiprocessing—distinct but similar)
  - Pipe (multiprocessing)
- Synchronization
  - Lock
  - Condition (i.e., condition variable)
  - Semaphore
  - ...

# Matrix Multiplication


## Basic Algorithm – Square Matrices



```
# SQUARE MATRICES
# C = A * B
# A, B, and C are n x n lists of lists.
# C is initialized to 0.
for i in range(n):
    for j in range(n):
        for k in range(n):
            C[i][j] += A[i][k] * B[k][j]
```

# Matrix Multiplication

## Basic Algorithm – Generalized



# GENERAL CASE

#  $C = A * B$

# A, B, and C are lists of lists.

# C is initialized to 0.

# A: m rows by n columns —  $m == \text{len}(A)$

# B: n rows by o columns —  $n == \text{len}(B) == \text{len}(A[0])$

# C: m rows by o columns —  $o == \text{len}(B[0])$

for i in range(len(A)):

    for j in range(len(B[0])):

        for k in range(len(B)):

$C[i][j] += A[i][k] * B[k][j]$

# Using One-Dimensional Arrays

## To Support Shared Memory

```
# C = A * B
# A, B, and C are multiprocessing Array objects.
# C is initialized to 0.
# A, B, and C all represent square n x n matrices.
for i in range(n):
    for j in range(n):
        for k in range(n):
            C[i * n + j] += A[i * n + k] * B[k * n + j]
```

```
# C = A * B
# A, B, and C are multiprocessing Array objects.
# C is initialized to 0.
# A: length m * n
# B: length n * o
# C: length m * o
for i in range(m):
    for j in range(o):
        for k in range(n):
            C[i * o + j] += A[i * n + k] * B[k * o + j]
```

# Dividing Up the Work

## Partitioning Data





- Are the data independent?
  - How do you know?
- Is any communication between processes necessary?
- How do you divide the input data up evenly among the workers?
  - 2?
  - 4?
  - 8?

Thing about it before we move on!






# Chunking

## Tiling



# Chunking

## Tiling




# Chunking


## Tiling




$p = 2^0$   
rows: 1  
columns: 1



$p = 2^1$   
rows: 2  
columns: 1



$p = 2^2$   
rows: 2  
columns: 2



$p = 2^3$   
rows: 4  
columns: 2




# Chunking


## Tiling




$p = 2^0$   
rows: 1  
columns: 1



$p = 2^1$   
rows: 2 (dim/2)  
columns: 1 (dim)



$p = 2^2$   
rows: 2 (dim/2)  
columns: 2 (dim/2)



$p = 2^3$   
rows: 4 (dim/2)  
columns: 2 (dim/4)




# Chunking

## Tiling




Matrix

$p = 2^0$   
rows: 1  
columns: 1




Matrix

$p = 2^1$   
rows: 2 (dim/2)  
columns: 1 (dim)  
vertical starts:  
0: 0 1: v\_chunk  
horizontal starts:  
0, 1: 0




Matrix

$p = 2^2$   
rows: 2 (dim/2)  
columns: 2 (dim/2)  
vertical starts:  
0, 1: 0  
2, 3: v\_chunk  
horizontal starts:  
0, 2: 0  
1, 3: h\_chunk



Matrix

$p = 2^3$   
rows: 4 (dim/2)  
columns: 2 (dim/4)  
vertical starts:  
0, 1: 0  
2, 3: v\_chunk  
4, 5: 2 \* v\_chunk  
6, 7: 3 \* v\_chunk  
horizontal starts:  
0, 2, 4, 6: 0  
1, 3, 5, 7: h\_chunk




# Chunking


## Tiling




$p = 2^0$   
rows: 1  
columns: 1



$p = 2^1$   
rows: 2 (dim/2)  
columns: 1 (dim)  
vertical starts:  
0:  $0 * v\_chunk$   
1:  $1 * v\_chunk$   
horizontal starts:  
0, 1:  $0 * h\_chunk$



$p = 2^2$   
rows: 2 (dim/2)  
columns: 2 (dim/2)  
vertical starts:  
0, 1:  $0 * v\_chunk$   
2, 3:  $1 * v\_chunk$   
horizontal starts:  
0, 2:  $0 * h\_chunk$   
1, 3:  $1 * h\_chunk$



$p = 2^3$   
rows: 4 (dim/2)  
columns: 2 (dim/4)  
vertical starts:  
0, 1:  $0 * v\_chunk$   
2, 3:  $1 * v\_chunk$   
4, 5:  $2 * v\_chunk$   
6, 7:  $3 * v\_chunk$   
horizontal starts:  
0, 2, 4, 6:  $0 * h\_chunk$   
1, 3, 5, 7:  $1 * h\_chunk$



Find the general pattern for vertical and horizontal starts (and ends) based on worker ID.

Test on  $p = 16 = 2^4$ .


# Chunking

## Tiling




Matrix

$p = 2^0$   
rows: 1  
columns: 1




Matrix

$p = 2^1$   
rows: 2 (dim/2)  
columns: 1 (dim)  
vertical starts:  
(pid // columns) \* v\_chunk  
horizontal starts:  
(pid % columns) \* h\_chunk



Matrix

$p = 2^2$   
rows: 2 (dim/2)  
columns: 2 (dim/2)  
vertical starts:  
(pid // columns) \* v\_chunk  
horizontal starts:  
(pid % columns) \* h\_chunk



Matrix

$p = 2^3$   
rows: 4 (dim/2)  
columns: 2 (dim/4)  
vertical starts:  
(pid // columns) \* v\_chunk  
horizontal starts:  
(pid % columns) \* h\_chunk

// = integer division

$1 // 2 = 0$

$2 // 2 = 1$


$3 // 2 = 1$

# Chunking


## Tiling




$p = 2^0$   
rows: 1  
columns: 1



$p = 2^1$   
rows: 2 ( $2^{q//2+1}$ )  
columns: 1 ( $2^{q//2}$ )  
vertical starts:  
(pid // columns) \* v\_chunk  
horizontal starts:  
(pid % columns) \* h\_chunk



$p = 2^2$   
rows: 2 ( $2^{q//2}$ )  
columns: 2 ( $2^{q//2}$ )  
vertical starts:  
(pid // columns) \* v\_chunk  
horizontal starts:  
(pid % columns) \* h\_chunk



$p = 2^3$   
rows: 4 ( $2^{q//2+1}$ )  
columns: 2 ( $2^{q//2}$ )  
vertical starts:  
(pid // columns) \* v\_chunk  
horizontal starts:  
(pid % columns) \* h\_chunk

// = integer division

$1 // 2 = 0$

$2 // 2 = 1$

$3 // 2 = 1$