

Machine Architectures and Parallel Computing

Week 4

Amittai Aviram – 29 September 2020

The Memory Cache

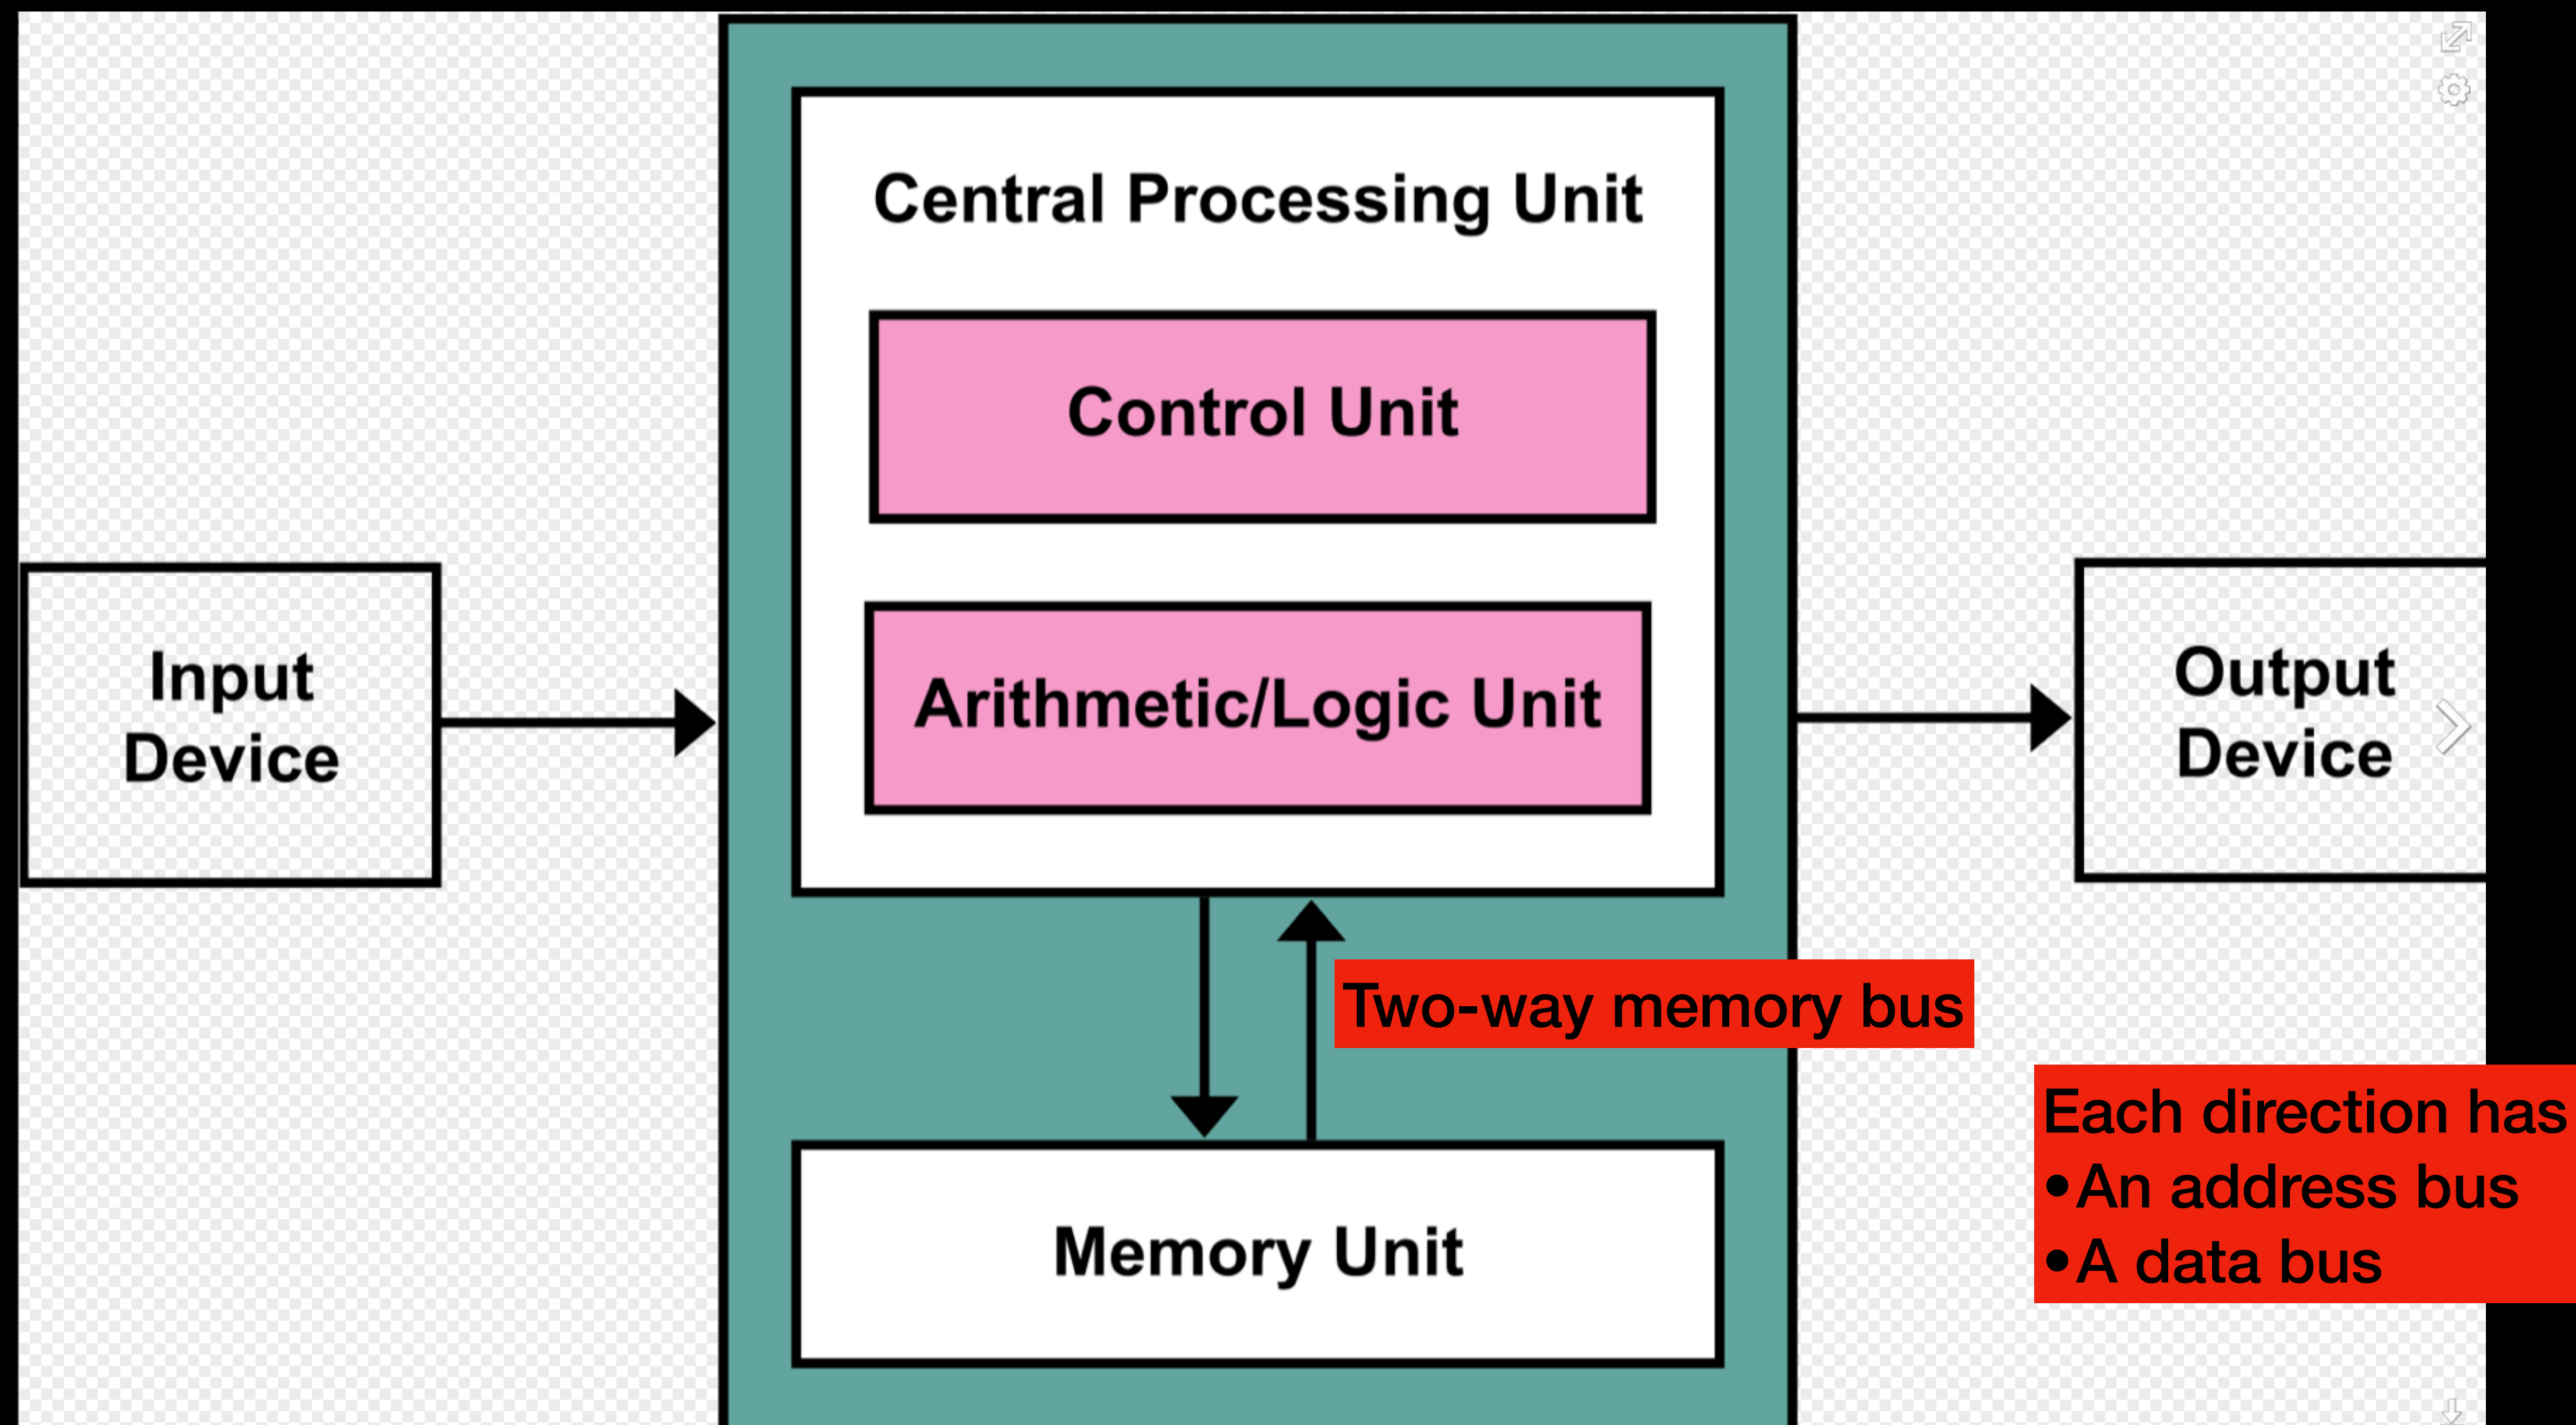
Overview

Why Are We Talking about Computer Architecture?

- Even in *shared memory parallelism*, we have to move data back and forth between processors and memory
- This is a hidden *communication cost*.
- This communication cost may be significant.
- We can mitigate it to improve the performance of shared memory parallelism.

The von Neumann Architecture

The Foundation of Digital Computers



The von Neumann Architecture

Published in 1945 by John von Neumann (1903 – 1957)



The von Neumann Bottleneck

Background

- A *bus* is a collection of wires or connections.
- The *memory bus* has as many connections as bits it carries.
 - Address bus—e.g., 64 connections for 64-bit addresses
 - Data bus—e.g., 64 connections for a 64-bit number
- *Fetch instruction*
 - CPU sets the bits for the address on the bus, along with a fetch signal
 - The memory unit sets the bits on the data bus corresponding to the value stored in memory
- *Store instruction*
 - The CPU sets the bits for the address on the bus, along with a store signal
 - It also sets the bits on the data bus corresponding to the value in the register to be stored

The von Neumann Bottleneck

The Problem

- Transmission of a signal across the bus is *much slower* than transmission between a register and the ALU in a CPU.
- Bus speed: 66 – 800 MHz
- Cache speed: > 1.8 GHz
- An instruction that requires a value from memory must cause the CPU to *block* until the value arrives.

The von Neumann Bottleneck

Example

- Processing Speed
 $8 \text{ cores} \times 3 \text{ MHz clock speed} \times 16 \text{ Flops per cycle} = 384 \text{ GFlops/s.}$
- Data Transfer Speed
 $800 \text{ MHz} \times 8 \text{ cores} \times 8 \text{ bytes per cycle} = 51.2 \text{ GB/s.}$
- Problem
Dot product of 2 vectors, length $2^{30} \rightarrow 2^{31}$ floating point operations (Flops) (2 GFlops).
- Data
 2^{31} 8-bit (double-precision) floating point numbers $\rightarrow 2^{34}$ bytes (16 GB)




The von Neumann Bottleneck

Example

- Computation time: $t_{\text{comp}} = \frac{2 \text{ GFlops}}{384 \text{ GFlops/s}} = 5.2 \text{ ms}$
- Data transfer time: $t_{\text{mem}} = \frac{16 \text{ GB}}{51.2 \text{ GB/s}} = 312.5 \text{ ms}$ **Memory bound**
- Assuming that memory access and computation are concurrent,
 $t_{\text{exec}} \geq \max(t_{\text{comp}}, t_{\text{data}}) = 312.5 \text{ ms}.$
- Performance: $\frac{2 \text{ GFlops}}{312 \text{ ms}} = 6.4 \text{ GFlops/s}$, and $\frac{6.4 \text{ GFlops/s}}{384 \text{ GFlops/s}} = 1.67 \%$

Solution

The Memory Cache (or Cache Memory)



The cache is a smaller intermediate container that holds copies of portions of the main memory that are likely to be used immediately. CPU-cache communication is *much faster* than CPU-memory communication.

Memory Cache

Example

- Data Storage Capacity
512 KB
- Data Transfer Speed
The same as for registers, so counted in with computation.
- Problem
Matrix multiplication with $n \times n$ matrices, $n = 128$.
For each element, n multiplication + n addition = $2n$ Flops.
 $2n \times n^2$ Flops = $2n^3$ Flops = $2 \cdot 128^3$ Flops = 2^{22} Flops = 4 GFlops
- Data
 $3 \times 128 \times 128 \times 8$ B = 3×128 KB = 384 KB < 512 KB.

```
for (size_t i = 0; i < n; ++i) {  
    for (size_t j = 0; j < n; ++j) {  
        for (size_t k = 0; k < n; ++k) {  
            W[i][j] += U[i][k] * V[k][j];  
        }  
    }  
}
```

Memory Cache

Example

- $t_{\text{comp}} = \frac{4 \text{ GFlops}}{384 \text{ GFlops/s}} = 10.4 \mu\text{s}$

Compute bound

- $t_{\text{mem}} = \frac{384 \text{ KB}}{51.2 \text{ GB/s}} = 7.5 \mu\text{s}$

We fetch data into the cache once at the beginning and store it back at the end, so computing and memory times do not overlap.

- $t_{\text{exec}} \geq t_{\text{comp}} + t_{\text{mem}} = 10.4 \mu\text{s} + 7.5 \mu\text{s} = 17.9 \mu\text{s}$

- Performance
 $\frac{4 \text{ GFlops}}{17.9 \mu\text{s}} = 223 \text{ GFlops/s}$, and $\frac{223 \text{ GFlops/s}}{384 \text{ GFlops/s}} = 58 \%$

Cache Algorithms

Common Solutions

- Cache access is much faster than main memory access, but a cache is much more expensive, so caches need to be small.
- Spatial locality — fetch data in *lines* — line size is, say, 64 bytes.
- Temporal locality — to make space for a new line, we evict the least-recently used one (LRU).
- Mapping of memory to cache addresses
 - Direct mapping — undermines LRU policy — higher miss rate.
 - Two-way (or n -way) mapping — a compromise.
 - Fully associative mapping — too expensive.

Cache-Based Optimization

Matrix Multiplication

- In the naïve algorithm, $B[k][j]$ skips from row to row (with the same column), so it causes a high miss rate.
- If you transpose B , you can get a significant improvement (by a factor of 2.2 in the Schmidt's example).

```
for (uint64_t i = 0; i < n; ++i) {  
    for (uint64_t j = 0; j < n; ++j) {  
        for (uint64_t k = 0; k < n; ++k) {  
            C[i][j] += A[i][k] * Bt[j][k]  
        }  
    }  
}
```

Cache Coherency

The Serial Case

- If the core needs to write to location x in main memory, a copy of which is stored in the cache, it would write to the cache so as to have the value immediately available.
- But this makes the cache line no longer a copy of main memory!
- Solutions:
 - *Write through*: always write *both* to the cache *and* to main memory. This is as slow as not having a cache.
 - *Write back*: mark the cache line as *dirty* and copy it to main memory when it is evicted. This introduces a period of cache incoherence.

Cache Coherency

The Parallel Case




FIGURE 3.5

Illustration of a CPU-chip with two cores and two levels of cache: a private L1-cache and a shared L2-cache.

See MESI Protocol:

https://en.wikipedia.org/wiki/MESI_protocol

In some cases, parallel cache interactions may cause a parallel program to run more slowly than its sequential equivalent.

Levels of Parallelism

Flynn's Taxonomy

Kinds of Parallel Architecture

- SISD — Single Instruction, Single Datum (not parallel)
- SIMD — Single Instruction, Multiple Data (vector instructions)
- Multiple Instructions, Multiple Data (independent processing elements)
- Multiple Instructions, Single Datum (pipelines)

Michael J. Flynn, Stanford, 1966

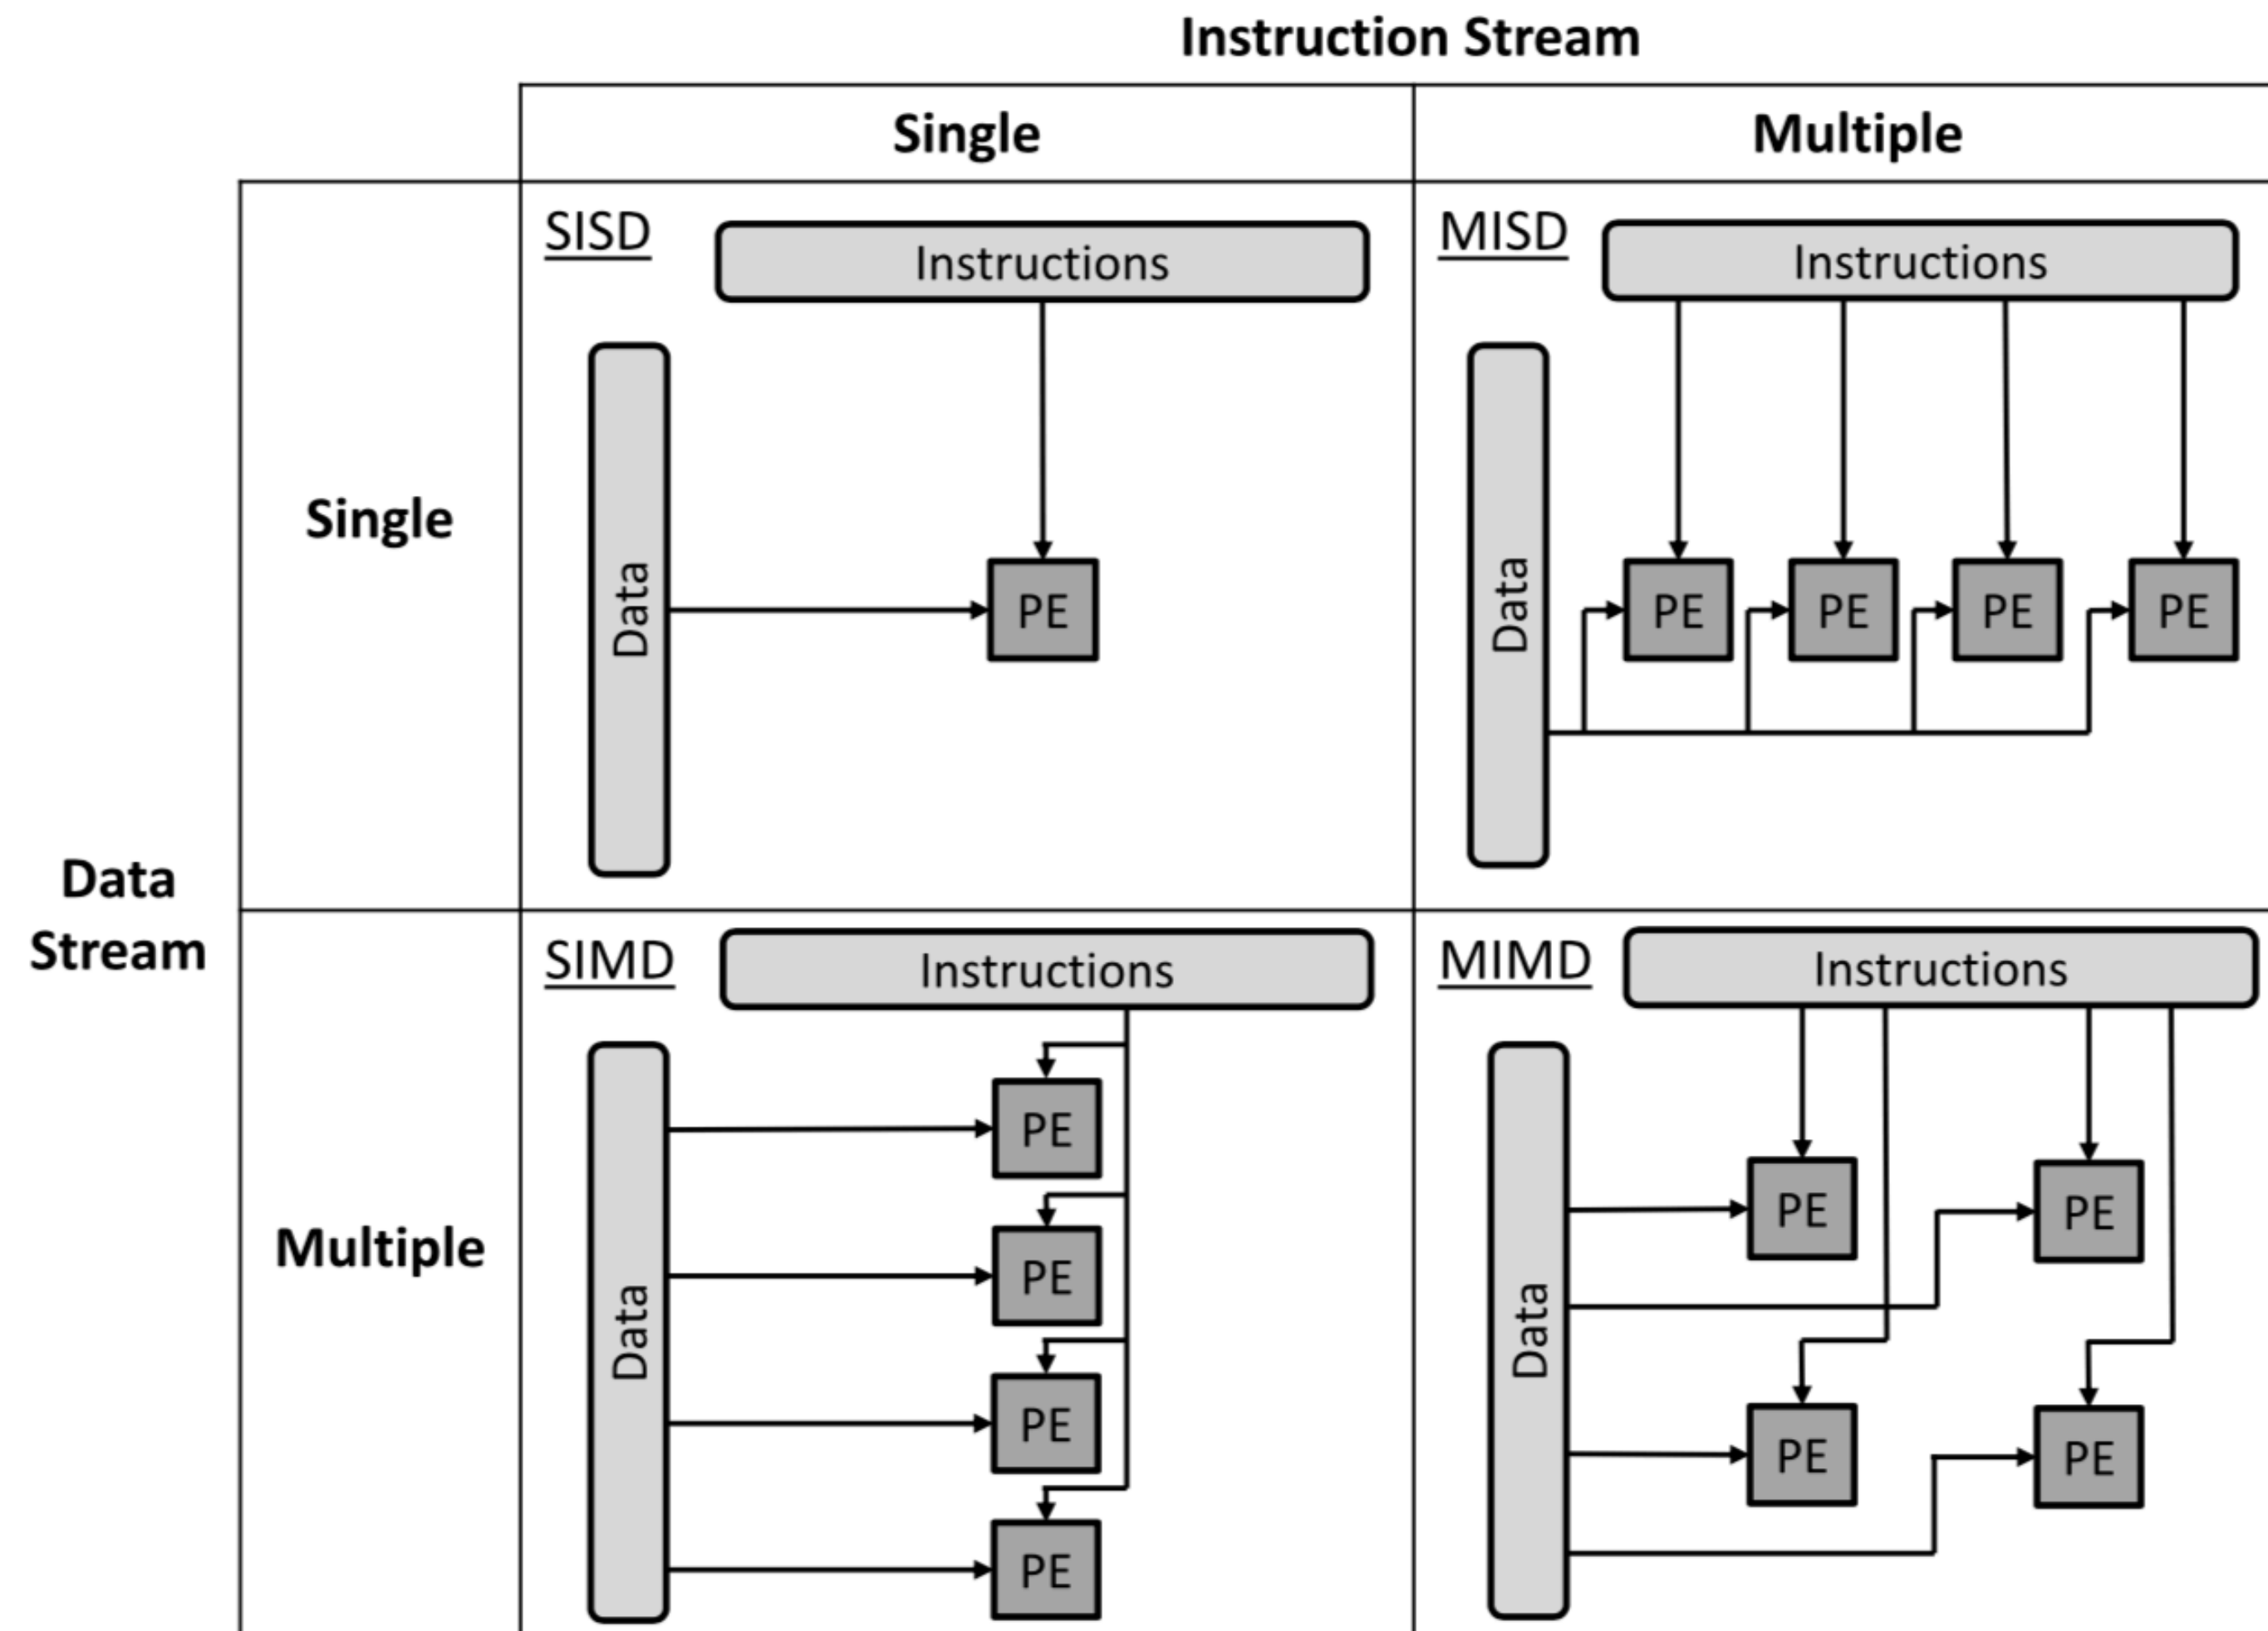


FIGURE 3.7

Flynn's taxonomy classifies computer architectures into four classes according to the number of instruction streams and data streams: SISD, MISD, SIMD, and MIMD.

Microarchitecture

Support for Flynn Classes of Parallelism

- MIMD — multiple cores.
- SIMD — vector registers, vector units, and vector instructions (“intrinsic”).
- MISD
 - Pipelining — simultaneous instruction fetching, instruction decoding, and data fetching.
 - Superscalar execution — pipelining of instructions in the ALU — requires out-of-order execution to avoid dependencies.

SIMD Instructions

Vectorized Registers and Operations

- Intel supports SIMD operations through *intrinsics*.
- For reference, see the [Intel Intrinsics Guide](#).

SIMD Optimizations

AoS versus SoA

- A SIMD operation over an array of structures (AoS) may be inefficient for vector hardware.
- Example
 - Array of length-3 arrays (vectors) of floating point numbers, representing an array of points in 3-dimensional space.
 - Each vector v_i must be *normalized*:
$$\hat{v}_i = \frac{v_i}{\|v_i\|} = \frac{x_i + y_i + z_i}{1/\sqrt{x_i^2 + y_i^2 + z_i^2}}$$
 - The triple (x_i, y_i, z_i) wastes space on vector hardware sized for a power of 2.
 - The sum of squares would have to work across vector hardware lanes, which is impossible, so it must be serialized.
- Translating the data to three arrays—a structure of arrays (SoA)—enables us to use vector hardware efficiently.

SIMD Optimization

AoS versus SoA

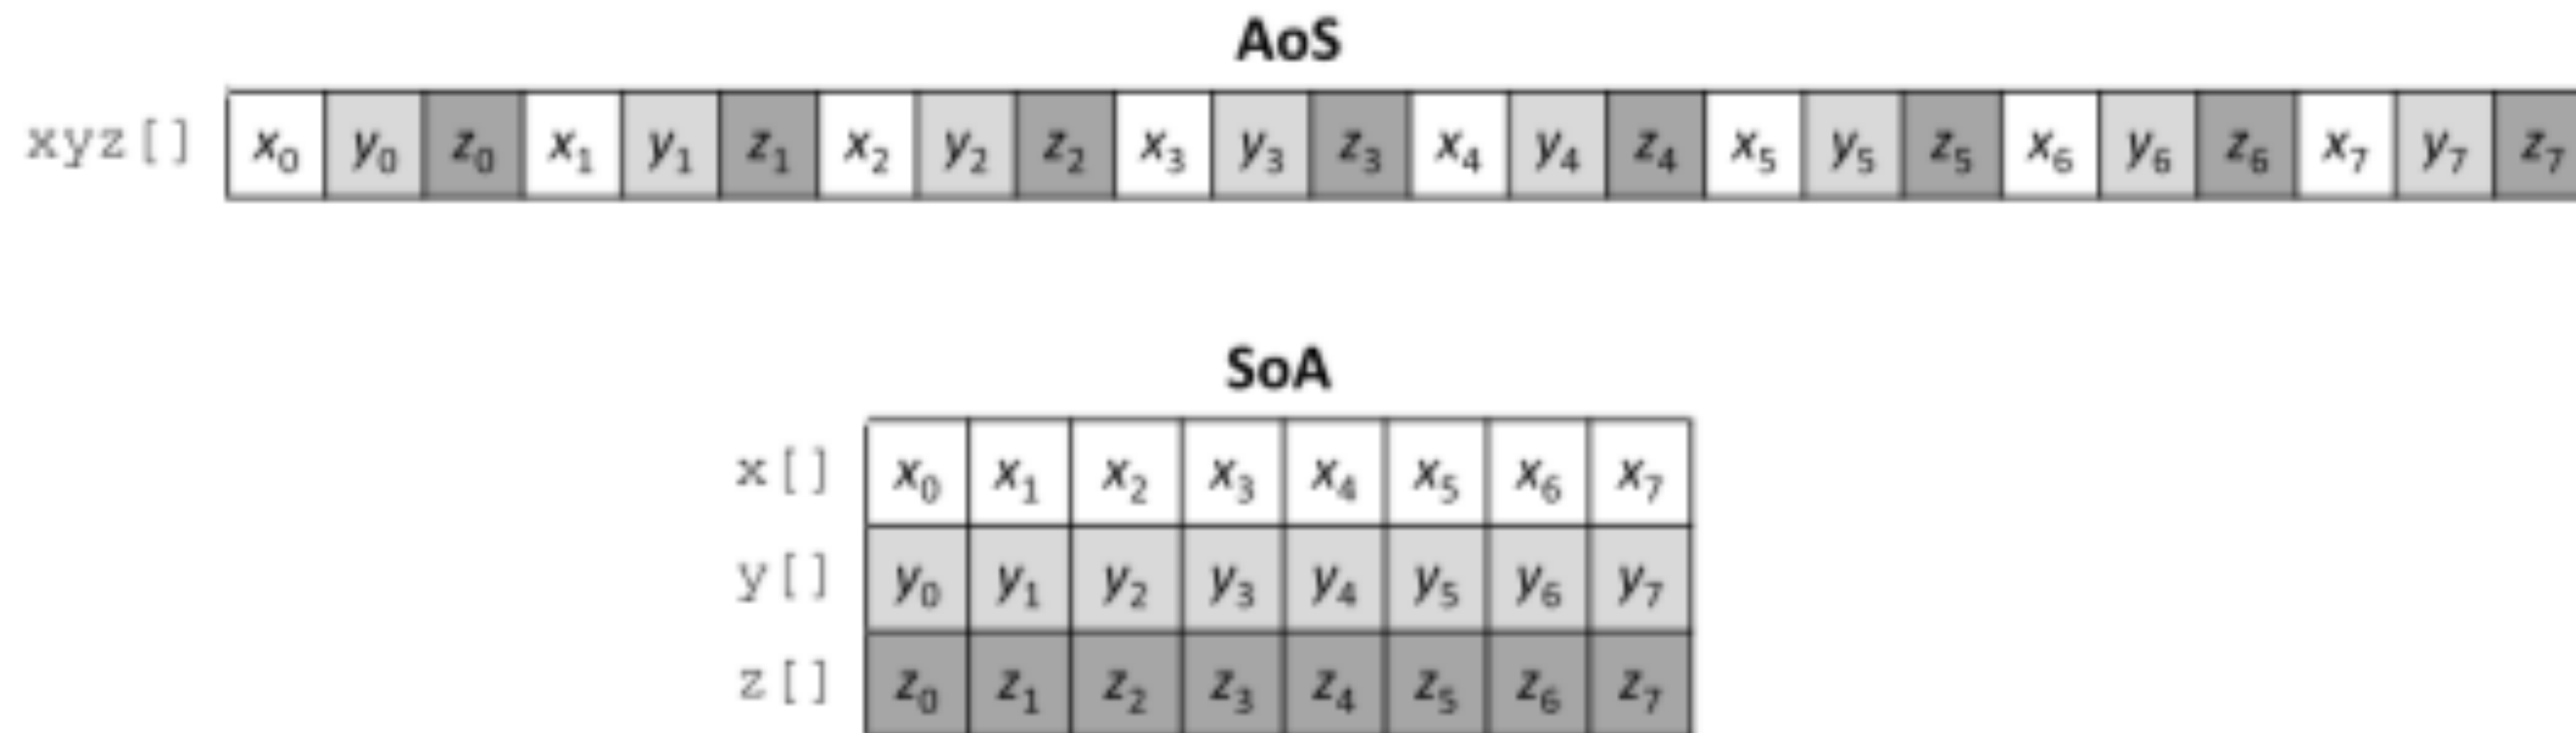


FIGURE 3.12

Comparison of the AoS and the SoA memory layout of a collection of eight 3D vectors:

Summary

Parallelism in the Machine Architecture Context

- Memory accesses can degrade the performance of parallel programs in a way comparable to the effect of excess communication between compute nodes.
- The memory cache mitigates this problem—if programs are refactored to take advantage of cache locality.
- Flynn's taxonomy of types of parallelism corresponds to various computer architecture features.
- Vector hardware and intrinsics support SIMD processing at the individual machine level.
- Programs can be refactored to benefit from vector hardware.

