

Parallel Programming in C++

Weeks 5 and 6

Amittai Aviram – 6 & 13 October 2020 – Boston University

Modern C++ Features

C++ Anonymous Functions

A Convenient Technique

- C++, Python, and other modern languages have syntax for *anonymous functions*.
- These enable you, e.g., to define a function *inline* when passing it as an argument to another function.
- Called “lambdas” after *lambda expressions* in the *lambda calculus* (Alonzo Church et al., 1932-36). (Sometimes called “closures,” but those also mean something else.)
- Syntax: [*capture symbols*] (*parameters*) { *body* }
- The *capture list* transfers in-scope symbols into the function’s environment.

Coding: anonymous functions.


C++ Smart Pointers

For Memory Allocation Control

- A smart pointer holds a (“raw”) pointer to allocated memory.
- It de-allocates the memory (calling `delete`) when it *goes out of scope*.
 - When it goes out of scope, the compiler inserts a call to its *destructor*.
 - The destructor calls `delete`.
- Types of smart pointer:
 - `unique_ptr` — allows only one scope to “own” the memory at a time.
 - `shared_ptr` — allows more than one scope to “own” the memory — uses a *reference counter* to call the destructor when the object goes out of *the last scope*.
 - `weak_ptr` — keeps track of a `shared_ptr` object.

L-Values and R-Values

C++ Subtleties



You can take the address:

```
int* p = &x;
```

You cannot take the address:

```
int* p = &3; // ERROR
```

You can assign:

```
x = 4;
```

You cannot assign:

```
3 = 4; // HAHA!
```

Variables and functions

Literals and temporaries

```
int x = 3 + y;
```

L-Value and R-Value References

- A *reference* works as an *alias* for an object.
- For “plain old data” (POD) objects, passing by l-value reference and passing by r-value reference are semantically identical.
- For complex objects,
 - Passing by l-value reference does not invoke a constructor.
 - Passing by r-value reference invokes a *move constructor*, which determines exact behavior.
- `std::move` casts an object as an r-value reference.

The `const` and `volatile` Modifiers

- `const` — makes an object *immutable*.
- `volatile` — make the compiler treat the object as if assignment to it had a side-effect (such as printing to **`stdout`**).

```
template <typename value_t, typename index_t>
value_t sequential_fibonacci(value_t n) {
    // initial conditions
    value_t a_0 = 0;
    value_t a_1 = 1;
    // iteratively compute the sequence
    for (index_t index = 0; index < n; ++index) {
        const value_t tmp = a_0;
        a_0 = a_1;
        a_1 += tmp;
    }
    return a_0;
}
```

C++ Concurrency

C++ Threads

Basics

- C++ threads are true kernel threads
 - Lightweight
 - In a shared address space
 - Assigned to processors as the latter become available
- Defined in the STL **<thread>** header
- Launched upon construction

Coding: “Hello, World!”

Returning a Value

Two Techniques

- A C++ thread has no mechanism for transferring a function's return value to the shared space using normal function return-value syntax.
- We use function parameters (“out-parameters”) on **void** functions instead:
 - A pointer type — holding the *address* in shared memory to which the function writes the return value in shared memory.
 - A *promise*, passed by *r-value reference*, to whose associated *promise* the function writes the return value.

Coding: sequential and parallel Fibonacci numbers.

Static Task Scheduling

Schemes to Divide Up the Work

- Block Distribution
Each thread works on a contiguous block of the shared data.
- Cyclic Distribution
Each thread id works on items $id, id + c, id + 2c, \dots$
This may worsen the problem of *false sharing*.
- Block-Cyclic Distribution
Each thread has every block of size c in each stride s .
This generalizes the other two schemes.
The smaller the block, the better the load balance (roughly).

Dynamic Scheduling

Determined at Runtime

- Each thread grabs a task, finishes it, and grabs the next task.
- Access to the source of tasks must be serialized
 - Lock
 - Condition variable

Synchronization Primitives

Used to Co-ordinate Threads and Access to Shared Data

- Mutual exclusion lock (**mutex**)
 - Use a *closure* to release the lock automatically — e.g., **lock_guard**.
- Condition Variable
 - Thread A waits
 - Thread B notifies
- Barrier
- Join

Coding: condition variables.

Lock-Free Synchronization

- Atomic types
 - These allow for an atomic change of value, including increment.
- Compare and swap (CAS)
 - Build atop an atomic type.
 - Used in a loop.
 - Say threads A and B want to add a subtotal to a sum. This requires
 - Read
 - Update
 - Write
 - By the time A reads and gets ready to write, B may have written, so the read is stale!
 - *A compares* its read value with the sum and writes only if these are equal. If not, it tries again.